Lab 2: Minnesota Tree Growth

AUTHOR   PUBLISHED
Samantha Nauman   February 21, 2025

Question 1: Describe the structure and what you see in the dataset?

In this dataset, it appears to be a 131,386 by 8 table of a tree growth analysis for 15 different tree species from 1897 to 2007. Within it, the columns consist of the treeID, standID, stand, year, species, age, inc (linear growth increment), and rad_ib (bark radius at breast height).

```
library(dplyr)
tree_dat <-read.csv("data/tree_dat.csv")
glimpse("tree_dat")
```

```
 chr "tree_dat"
```

Question 2: How many records have been made in stand 1?

In stand 1, there have been 979 records made.

```
tree_dat %>%
   filter(standID == 1) %>%
   tally()
```

```
    n
1 979
```

Question 3: How many records of the Abies balsamea and Pinus strobus species have been made?

For the Abies balsamea, there have been 13,033 records made. For the Pinus strobus, there have been 4,188 records made.

```
tree_dat %>%
   filter(species %in% c("ABBA", "PIST")) %>%
   count(species)
```

```
  species      n
1    ABBA 13033
2    PIST  4188
```

Question 4: How many trees are older than 200 years old in the last year of the dataset?

In the last year of the tree_dat dataset, there are 7 trees older than 200 years.

```
last_year <- max(tree_dat$year, na.rm = TRUE) #finds last yr in the data set

tree_dat %>%
  filter(year == last_year, age > 200) %>%
  tally() #counts trees older than 200 yrs in the last yr
```

```
    n
1 7
```

Question 5: What is the oldest tree in the dataset found using slice_max?

The oldest tree in the dataset is the Pinus resinosa, which is 296 years old.

```
tree_dat %>%
  slice_max(order_by = age, n=1)
```

```
   treeID standID stand year species age  inc rad_ib
1      24       2    A2 2007    PIRE 269 0.37 308.84
```

Question 6: Find the oldest 5 trees recorded in 2001.

The 5 oldest trees recorded in 2001 are two Pinus resinosa (263, 259) and three Fraxinus nigra (212, 206, 206).

```
tree_dat %>%
  filter(year == 2001) %>%
  slice_max(order_by = age, n = 5)
```

```
   treeID standID stand year species age   inc   rad_ib
1      24       2    A2 2001    PIRE 263 0.210  306.880
2      25       2    A2 2001    PIRE 259 0.280  156.210
3    1595      24    F1 2001    FRNI 212 0.579  156.267
4    1598      24    F1 2001    FRNI 206 0.394  130.251
5    1712      26    F3 2001    FRNI 206 0.168  154.354
```

Question 7: Using slice_sample, how many trees are in a 30% sample of those recorded in 2002?

In a 30% sample of recorded trees in 2002, there are 687 trees. (did not show output since it was a large table)

```
tree_dat %>%
  filter(year == 2002) %>%
  slice_sample(prop = 0.3)
```

Question 8: Filter all trees in stand 5 in 2007. Sort this subset by *descending* radius at breast height (rad_ib) and use slice_head() to get the top three trees. Report the tree IDs.

In descending order of tree radius at breast height in stand 5 in 2007, the trees are #128 (Pinus strobus), #157 (Pinus resinosa), and #135 (Picea mariana).

```
tree_dat %>%
   filter(standID == 5, year == 2007) %>%
   arrange(desc(rad_ib)) %>%
   slice_head(n = 3)
```

	treeID	standID	stand	year	species	age	inc	rad_ib
1	128	5	A6	2007	PIST	82	0.885	238.8850
2	157	5	A6	2007	PIRE	85	0.900	217.8700
3	135	5	A6	2007	PIMA	84	0.110	210.1874

Question 9: Reduce your full data.frame to [treeID, stand, year, and radius at breast height]. Filter to only those in stand 3 with records from 2007, and use slice_min to pull the smallest three trees measured that year.

The smallest three trees in stand 3 measured in 2007 in descending order are #50, #56, and #36.

```
tree_dat %>%
   select(treeID, standID, year, rad_ib) %>%
   filter(standID == 3, year == 2007) %>%
   slice_min(rad_ib, n = 3)
```

	treeID	standID	year	rad_ib
1	50	3	2007	47.396
2	56	3	2007	48.440
3	36	3	2007	54.925

Question 10: Use select to *remove* the stand column. Use glimpse to show the dataset.

```
tree_dat %>%
   select(-stand) %>%
   glimpse()
```

```
Rows: 131,386
Columns: 7
$ treeID  <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,…
$ standID <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,…
$ year    <int> 1960, 1961, 1962, 1963, 1964, 1965, 1966, 1967, 1968, 1969, 19…
$ species <chr> "ABBA", "ABBA", "ABBA", "ABBA", "ABBA", "ABBA", "ABBA", "ABBA"…
$ age     <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18,…
```

```
$ inc      <dbl> 0.930, 0.950, 0.985, 0.985, 0.715, 0.840, 0.685, 0.940, 1.165,…
$ rad_ib   <dbl> 10.78145, 11.73145, 12.71645, 13.70145, 14.41645, 15.25645, 15…
```

Question 11: Look at the help document for dplyr::select and examine the "Overview of selection features". Identify an option (there are multiple) that would help select all columns with the string "ID" in the name. Using glimpse to view the remaining dataset.

```
?dplyr::select
tree_dat %>%
    select(contains("ID")) %>%
    glimpse()
```

```
Rows: 131,386
Columns: 2
$ treeID   <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,…
$ standID  <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,…
```

Question 12: Find a selection pattern that captures all columns with either 'ID' or 'stand' in the name. Use glimpse to verify the selection.

```
tree_dat %>%
    select(contains("ID"), contains("stand")) %>%
    glimpse()
```

```
Rows: 131,386
Columns: 3
$ treeID   <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,…
$ standID  <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,…
$ stand    <chr> "A1", "A1", "A1", "A1", "A1", "A1", "A1", "A1", "A1", "A1", "A…
```

Question 13: Looking back at the data dictionary, rename rad_ib and inc to include _[unit] in the name. Unlike earlier options, be sure that this renaming is permanent, and stays with your data.frame (e.g. <-). Use glimpse to view your new data.frame.

```
tree_dat <- tree_dat %>%
    rename(rad_ib_mm = rad_ib, inc_mm = inc) %>%
    glimpse()
```

```
Rows: 131,386
Columns: 8
$ treeID    <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, …
$ standID   <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, …
$ stand     <chr> "A1", "A1", "A1", "A1", "A1", "A1", "A1", "A1", "A1", "A1", …
$ year      <int> 1960, 1961, 1962, 1963, 1964, 1965, 1966, 1967, 1968, 1969, …
```

```
$ species   <chr> "ABBA", "ABBA", "ABBA", "ABBA", "ABBA", "ABBA", "ABBA", "ABB…
$ age       <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 1…
$ inc_mm    <dbl> 0.930, 0.950, 0.985, 0.985, 0.715, 0.840, 0.685, 0.940, 1.16…
$ rad_ib_mm <dbl> 10.78145, 11.73145, 12.71645, 13.70145, 14.41645, 15.25645, …
```

Question 14: A key measurement in forestry in "basal area column". The metric is computed with the formula:

$BA(m2) = 0.00007854 \cdot DBH^2$ Where DBH is the diameter at breast height (cm). Use mutate to compute DBH in centimeters, and BA in m2 (HINT: Make sure rad_ib is in cm prior to computing the diameter!). What is the mean BA_m2 of the species POTR in 2007?

The mean BA_m2 of the species POTR in 2007 is 0.03696619 m^2.

```
tree_dat %>%
  mutate(
    rad_ib_cm = rad_ib_mm/10,
    DBH = rad_ib_cm*2,
    BA_m2 = 0.00007854*DBH^2) %>%
  filter(species == "POTR", year == 2007) %>%
  summarize(mean_BA_m2 = mean(BA_m2, na.rm = TRUE)) %>%
  print()
```

```
  mean_BA_m2
1 0.03696619
```

Question 15: Lets say for the sake of our study, trees are not established until they are 5 years of age. Use if_else to add a boolean column to our dataset called established that is TRUE if the age is greater then 5 and FALSE if less then or equal to five. Once added, use count (see ?count) to determine how many records are from estabilshed trees?

From established trees, there are 122,503 records.

```
tree_dat <- tree_dat %>%
  mutate(established = if_else(age > 5, TRUE, FALSE))
tree_dat %>%
count(established)
```

```
  established      n
1       FALSE   8883
2        TRUE 122503
```

Question 16: Use mutate and case_when to add a new column to your data.frame that classifies each tree into the proper DBH_class. Once done, limit your dataset to the year 2007 and report the number of each class with count.

In the year 2007, there are 0 seedlings, 252 saplings, 1963 poles, and 76 sawlogs.

```r
tree_dat <- tree_dat %>%
  mutate(rad_ib_cm = rad_ib_mm/10,
    DBH = rad_ib_cm*2) %>%
  mutate(DBH_class = case_when(
    DBH >= 0 & DBH < 2.5 ~ "seedling",
    DBH >= 2.5 & DBH < 10 ~ "sapling",
    DBH >= 10 & DBH < 30 ~ "pole",
    DBH >= 30 ~ "sawlog",
    TRUE ~ NA_character_
  ))
tree_dat %>%
  filter(year == 2007) %>%
  count(DBH_class) %>%
  print()
```

```
  DBH_class    n
1      pole 1963
2   sapling  252
3    sawlog   76
```

Question 17: Compute the mean DBH (in cm) and standard deviation of DBH (in cm) for all trees in 2007. Explain the values you found and their statistical meaning.

For all trees in 2007, the mean DBH (cm) is 16.09351, and the standard deviation of DBH (cm) is 6.138643.

```r
tree_dat %>%
  filter(year == 2007) %>%
  summarize(mean_DBH = mean(DBH, na.rm = TRUE),
            std_DBH = sd(DBH, na.rm = TRUE)) %>%
  print()
```

```
  mean_DBH   std_DBH
1 16.09351 6.138643
```

Question 18: Compute the *per species* mean tree age using only those ages recorded in 2003. Identify the three species with the oldest mean age.

The three species with the oldest mean tree age in descending order is THOC (Thuja occidentalis) at 126.63830 years, FRNI (Fraxinus nigra) at 83.08333 years, and PIST (Pinus strobus) at 73.28571 years old.

```r
tree_dat %>%
  filter(year == 2003) %>%
```

```
  group_by(species) %>%
  summarize(mean_tree_age = mean(age, na.rm = TRUE)) %>%
  arrange(desc(mean_tree_age)) %>%
  slice(1:3) %>%
  print()
```

```
# A tibble: 3 × 2
  species mean_tree_age
  <chr>           <dbl>
1 THOC            127.
2 FRNI             83.1
3 PIST             73.3
```

Question 19: In a single summarize call, find the number of unique years with records in the data set along with the first and last year recorded?

There are 111 unique years with records in the data set, with the first year recorded in 1897 and the last year recorded in 2007.

```
tree_dat %>%
  summarize(unique_years = n_distinct(year, na.rm = TRUE),
            first_year = min(year, na.rm = TRUE),
            last_year = max(year, na.rm = TRUE)) %>%
  print()
```

```
  unique_years first_year last_year
1          111       1897      2007
```

Question 20: Determine the stands with the largest number of unique years recorded. Report all stands with largest (or tied with the largest) temporal record.

The stands all tied with the largest temporal record at 111 unique years is stands A1, D1, D2, D3, and F1.

```
tree_dat %>%
  group_by(stand) %>%
  summarize(unique_years = n_distinct(year, na.rm = TRUE)) %>%
  ungroup() %>%
  filter(unique_years == max(unique_years)) %>%
  print()
```

```
# A tibble: 5 × 2
  stand unique_years
  <chr>        <int>
1 A1             111
2 D1             111
3 D2             111
```

```
4 D3                 111
5 F1                 111
```

Final Question: Use a combination of dplyr verbs to compute these values (per year growth made by each tree, total growth record, avg and std across the species) and report the 3 species with the fastest growth, and the 3 species with the slowest growth. (** You will need to use either lag() or diff() in your compuation. You can learn more about each in the Help pages). Lastly, find and include an image of the fastest growing species. Add the image to your images directory.

Top 3 Fastest Growing Species: PIRE (Pinus resinosa), POTR (Populus tremuloides), PIBA (Pinus banksiana)

Top 3 Slowest Growing Species: LALA (Larix laricina), THOC (Thuja occidentalis), QURU (Quercus rubra)

```r
temp <- tree_dat %>%
  group_by(treeID) %>%
  arrange(year) %>%
  mutate(growth = DBH - lag(DBH),
    total_years = n()) %>%
  filter(!is.na(growth), total_years >10) %>%
  ungroup() %>%
  group_by(species) %>%
  summarize(mean_growth = mean(growth))

 fastest <- slice_max(temp, mean_growth, n = 3)
 slowest <- slice_min(temp, mean_growth, n = 3)

 print(fastest)
```

```
# A tibble: 3 × 2
  species mean_growth
  <chr>        <dbl>
1 PIRE         0.358
2 POTR         0.331
3 PIBA         0.326
```

```r
 print(slowest)
```

```
# A tibble: 3 × 2
  species mean_growth
  <chr>        <dbl>
1 LALA         0.150
2 THOC         0.153
3 QURU         0.168
```

Fastest Growing Species: Pinus resinosa